

CIS 457 - Data Communications

Nathan Bowman

Images taken from Kurose and Ross book

TCP Congestion Control -- Policy

Final question for TCP congestion control is policy: how much do we lower or raise rate based on ACKs and losses?

Best way depends on network and traffic characteristics

Tricky balancing act, because too high of rate causes congestion, but too low of rate wastes bandwidth

Protocol must be distributed -- individual TCP senders control only their own behavior and know only what they directly observe from network

To further complicate matters, conditions in network
may change over time

TCP constantly probing to attempt to determine
highest rate it can send at right now

Any time ACK received, sender increases transmission
rate

Keeps increasing until loss event, at which point it
decreases

Cyclical process: after decrease, sender goes right back to increasing rate

By always "pushing the limit", sender can keep track of changes in network

Algorithm for determining rate split into three parts or phases

- slow start
- congestion avoidance
- fast recovery

Slow start and congestion avoidance are mandatory parts of protocol, but fast recovery is optional (recommended)

Slow start

Sender always begins in slow start


In slow start, cwnd initially set to just 1 MSS

Transmission rate therefore MSS/RTT

This lives up to the "slow" part -- sender wants to increase rate quickly

Sender increases window size by 1 MSS per new ACK

Does not sound fast, but consider that number of ACKs increases each time window size increases



- Send 1 MSS; wait RTT; ACK
- Send 2 MSS; wait RTT; 2 ACKs
- Send 4 MSS; wait RTT; 4 ACKs
- Send 8 MSS; wait RTT; 8 ACKs

Window size doubles each RTT -- this is exponential growth

Exponential growth can only continue so long

Three ways to leave slow start:

- packet loss (timeout)
- hit threshold based on most recent packet loss
- triple duplicate ACKs

When packet loss occurs, sender sets **slow start threshold** (`ssthresh`) to $cwnd/2$

Sender then resets `cwnd` to 1 and starts over, now with new stopping point `ssthresh`

If no loss events occur before `ssthresh` reached,
sender still ceases to double rate after `ssthresh`

Doing so would cause rate to jump to where last packet
loss occurred

Instead, sender leaves slow start stage and enters
congestion avoidance

Third and final way to trigger end of slow start is receipt of triple duplicate ACKs

When this occurs, sender performs fast retransmit (as discussed in previous lecture) and enters fast recovery state

Congestion avoidance

In congestion avoidance stage, sender increases cwnd by 1 MSS per RTT (rather than per ACK)

How can sender know RTT without clocking it?

Still a matter of counting ACKs

At any time, sender has roughly cwnd/MSS packets in network

Should receive roughly cwnd/MSS ACKs per RTT

So, increase by $\text{MSS} * (\text{MSS}/\text{cwnd})$ bytes per new ACK

For example, if MSS is 1460 and cwnd is 14600, each ACK will increase cwnd by

$$1460 * (1460 / 14600) \\ = 1/10 * 1460$$

cwnd was 10 packets, and increases by 1/10 packet per ACK, which is roughly 1 packet per RTT

Congestion avoidance is slower, safer linear increase

No upper threshold -- continues increasing until loss event

When timeout occurs, do exactly same thing as slow start would:

Set `ssthresh` to `cwnd/2`

Set `cwnd` to 1

Restart in slow start

In case of triple duplicate ACK, also does same thing as slow start would, though we have not specified that yet

This scenario is different from timeout because it tells sender several packets have still made it through, so perhaps congestion not so bad

Sender performs fast retransmit

Sets $ssthresh$ to $cwnd/2$

Rather than start over, set $cwnd$ to $cwnd/2 + 3$
(because of 3 ACKs)

Go to fast recovery state

Fast recovery

This is temporary state to deal with loss event in network that may not be especially congested

Increase cwnd by 1 MSS for each duplicate ACK of packet that put us into this state

If timeout occurs, behaves exactly the same as other two stages:

- set `ssthresh` appropriately,
- set `cwnd` to 1 MSS, and
- enter slow start

This state was based on lost packet, and ends when
ACK for that packet arrives


At that point, sender will:

- set cwnd to ssthresh
- enter congestion avoidance

As mentioned, fast recovery is optional

Older TCP standard, TCP Tahoe, did not have fast
recovery

Newer TCP Reno does



Like many protocols, congestion control in TCP can be described succinctly by FSM

